

## Experiment-09

### **Cloud Monitoring: Discover how Cloud Monitoring help in tracking and analyzing the performance and health of cloud resources?**

**Google Cloud Monitoring** (formerly **Stackdriver Monitoring**) provides robust monitoring, alerting, and dashboard capabilities to track and analyze the performance and health of resources in your Google Cloud environment. It helps you ensure that your cloud infrastructure and applications are functioning efficiently and can scale as needed.

#### **Key Features of Google Cloud Monitoring**

##### **Performance Tracking**

- **Resource Metrics** – Track metrics for **Compute Engine**, **Kubernetes Engine**, **Cloud Functions**, **Cloud SQL**, and other services to assess resource usage (e.g., CPU, memory, disk, network usage).
- **Custom Metrics** – You can define your own custom metrics for specific applications to monitor application-level health or performance.
- **Real-Time Metrics** – Get near real-time data updates on how your infrastructure is performing.

##### **Health Monitoring**

- **Health Checks** – Cloud Monitoring can check the health of your **Compute Engine** instances, **App Engine**, or any other services you configure, ensuring that everything is running smoothly.

- **Uptime Checks** – Automate the monitoring of service availability across regions and ensure that downtime is minimized.

## Alerting and Notifications

- **Smart Alerts** – Create alerts based on specific thresholds for metrics such as CPU utilization, memory usage, or response time.
- **Notification Channels** – Alerts can be routed through email, SMS, Slack, or other communication platforms, ensuring that you are immediately notified about any issues.
- **Escalation Policies** – Implement escalation policies so that alerts are sent to the right teams if an issue persists.

## Dashboards and Visualization

- **Custom Dashboards** – Create visual dashboards that provide a snapshot of the health and performance of your Google Cloud services, apps, and infrastructure.
- **Predefined Dashboards** – Use built-in dashboards for common Google Cloud resources like **Compute Engine**, **Kubernetes Engine**, and **Cloud Pub/Sub**.

## Distributed Tracing and Logging

- **Distributed Tracing** – Track the performance of services that interact with one another, pinpointing latency or bottlenecks in distributed systems, microservices, and serverless architectures.
- **Log Analysis** – Cloud Monitoring integrates with **Cloud Logging** to collect logs from your applications and infrastructure. It helps you quickly diagnose issues by correlating logs with performance metrics.

## **Integration with Cloud Services**

- **Cloud Monitoring for Kubernetes** – Keep track of the health and performance of your Kubernetes clusters, containers, and pods.
- **Integration with Google Cloud Services** – Seamlessly integrates with **Compute Engine**, **Google Kubernetes Engine**, **Cloud Functions**, and **Cloud Run** to give you full visibility into your infrastructure.

## **Use Cloud Monitoring**

### **Step 1: Enable Cloud Monitoring API**

1. Open the **Google Cloud Console** → **Navigation Menu (Ξ)** → **APIs & Services** → **Library**.
2. Search for **Cloud Monitoring API** and click **Enable**.

### **Step 2: Set Up Monitoring for Your Resources**

1. Go to **Navigation Menu (Ξ)** → **Monitoring** → **Dashboards**.
2. Click on **Create Dashboard** to start building your custom dashboard.
3. Select **Metrics** from the dropdown and choose the service you want to monitor (e.g., **Compute Engine**, **Cloud Storage**, **Cloud Functions**).
4. Add the required metrics to your dashboard and adjust visualizations like **line charts**, **heat maps**, or **bar charts**.

## **Step 3: Set Up Alerts**

1. Go to **Navigation Menu ( $\equiv$ )** → **Monitoring** → **Alerting**.
2. Click **Create Policy**.
3. Choose the **condition** (e.g., CPU usage  $> 80\%$ ).
4. Select the **notification channels** (e.g., email, Slack, SMS) where the alert should be sent.
5. Set up **escalation policies** to ensure the right team is notified.
6. Click **Create** to finish the alert policy.

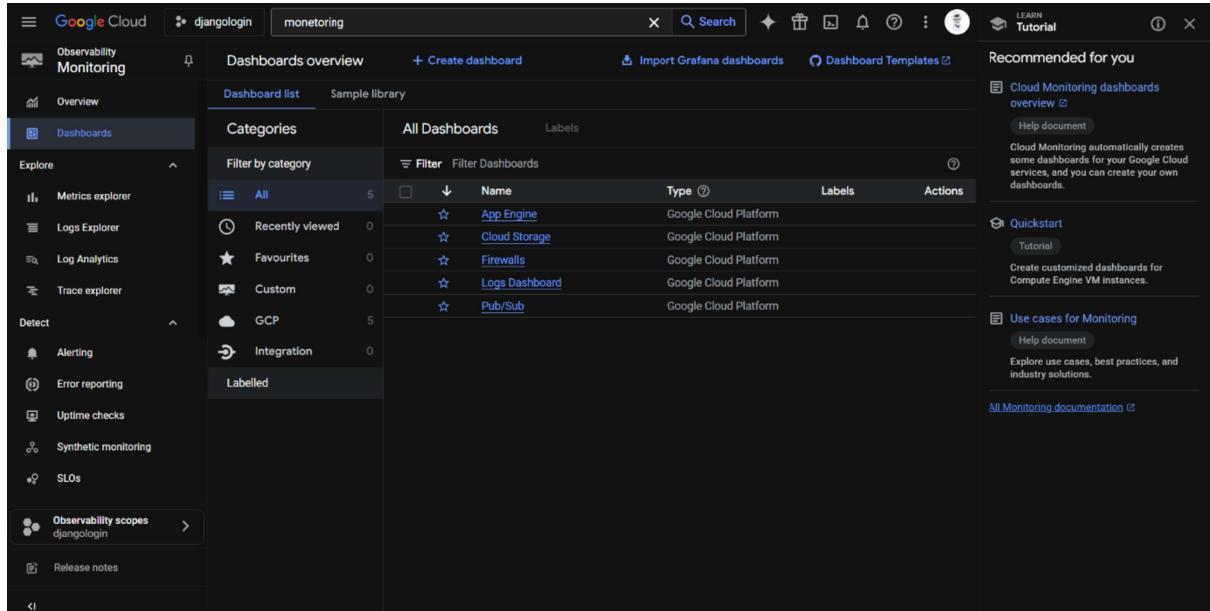
## **Step 4: Review Logs**

1. Go to **Navigation Menu ( $\equiv$ )** → **Logging**.
2. Choose the **resource type** (e.g., VM instances, Kubernetes Engine) and define your log filter.
3. Use **Log Explorer** to search for specific logs, such as **error messages** or **performance warnings**, and correlate them with metrics in Cloud Monitoring.

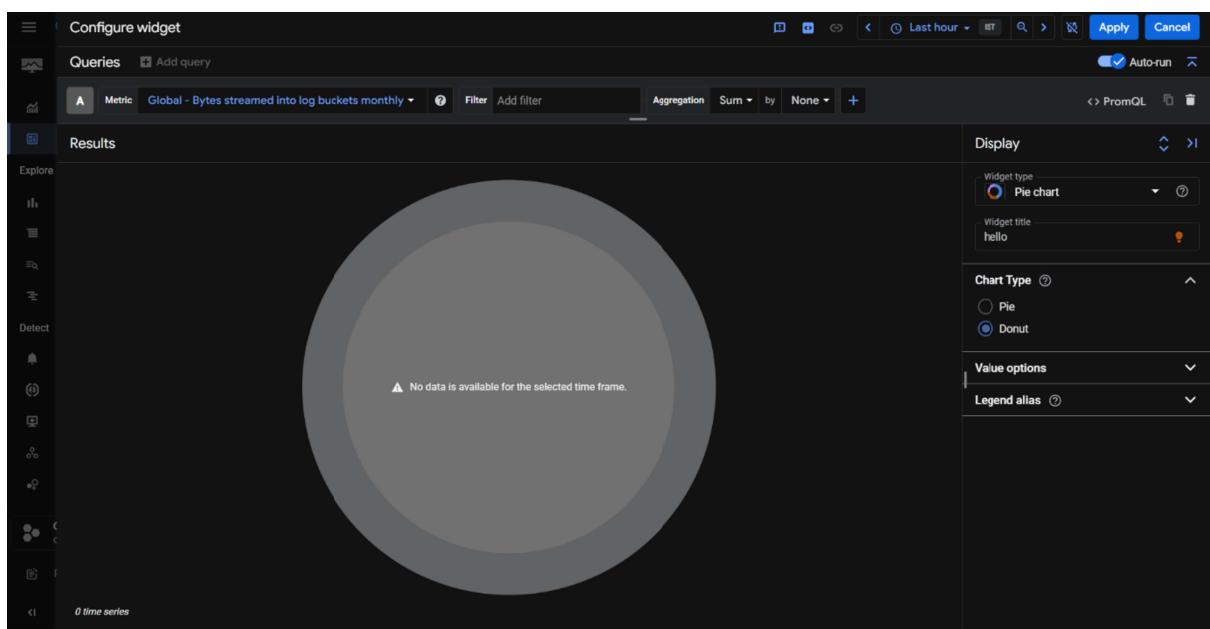
## **Step 5: Set Up Distributed Tracing (Optional)**

1. Open **Cloud Trace** from the **Navigation Menu**.
2. Enable **trace collection** in your services (e.g., by using the **Cloud Trace SDK** for your app).
3. View trace data to identify latency or bottlenecks in your system.

# Output



The screenshot shows the Google Cloud Monitoring Dashboards overview page. The left sidebar includes sections for Observability Monitoring (Overview, Dashboards), Explore (Metrics explorer, Logs Explorer, Log Analytics, Trace explorer), Detect (Alerting, Error reporting, Uptime checks, Synthetic monitoring, SLOs), and Observability scopes (djangologin). The main area displays a dashboard list with categories: Recently viewed (0), Favourites (0), Custom (0), GCP (5), and Pub/Sub (0). A 'Labelled' filter is applied. The dashboard list table has columns for Name, Type, Labels, and Actions. Recommended sections include 'Cloud Monitoring dashboards overview', 'Quickstart', and 'Use cases for Monitoring'.



The screenshot shows the 'Configure widget' screen. The top navigation bar includes 'Configure widget', 'Last hour', 'Apply', and 'Cancel'. The main area is titled 'Results' and shows a large gray circle with the message 'No data is available for the selected time frame.' Below this, it says '0 time series'. The configuration sidebar on the right is titled 'Display' and contains settings for 'Widget type' (Pie chart selected), 'Widget title' (hello), 'Chart Type' (Donut selected), 'Value options', and 'Legend alias'.